

# CS561 HW14

December 9, 2023

## 1 Problem 1

We define the entropy of continuous random variable  $X$  as follows:

$$H(X) = \mathbb{E}_X[-\log(p(x))] = \int_{\mathcal{X}} -p(x) \log(p(x)) dx$$

- a) Let  $\mathbf{x} \in \mathbb{R}^n$  be a random vector,  $\mu$  be a constant vector, and define  $\mathbf{y} = \mathbf{x} + \mu$ . We can compute the entropy of  $\mathbf{y}$  as follows:

$$\begin{aligned} H(\mathbf{y}) &= \mathbb{E}_{\mathbf{x}}[-\log(p(\mathbf{x} + \mu))] \\ &= \int_{\mathbb{R}^n} -p(\mathbf{x} + \mu) \log(p(\mathbf{x} + \mu)) d\mathbf{x} \\ &= \int_{\mathbb{R}^n} -p(\mathbf{x} + \mu) \log(p(\mathbf{x} + \mu)) d(\mathbf{x} + \mu) \\ &= \int_{\mathbb{R}^n} -p(\mathbf{u}) \log(p(\mathbf{u})) d\mathbf{u} \\ &= H(\mathbf{x}) \end{aligned}$$

- b) Let  $\mathbf{x} \in \mathcal{N}(0, \Sigma)$ , we can compute the entropy of  $\mathbf{x}$  as follows:

$$\begin{aligned} H(\mathbf{x}) &= -\mathbb{E}[\log(p(\mathbf{x}))] \\ &= \frac{1}{2} \mathbb{E} [\log((2\pi)^n |\Sigma|) + \mathbf{x} \Sigma^{-1} \mathbf{x}^T] \\ &= \frac{\log((2\pi)^n |\Sigma|)}{2} + \frac{1}{2} \mathbb{E} [\mathbf{x} \Sigma^{-1} \mathbf{x}^T] \\ &= \frac{n \log(\pi)}{2} + \frac{|\Sigma|}{2} + \frac{1}{2} \mathbb{E} [\text{tr}(\mathbf{x} \Sigma^{-1} \mathbf{x}^T)] \\ &= \frac{n \log(\pi)}{2} + \frac{|\Sigma|}{2} + \frac{1}{2} \mathbb{E} [\text{tr}(\Sigma^{-1} \mathbf{x}^T \mathbf{x})] \\ &= \frac{n \log(\pi)}{2} + \frac{|\Sigma|}{2} + \frac{1}{2} \text{tr}(\Sigma^{-1} \mathbb{E} [\mathbf{x}^T \mathbf{x}]) \\ &= \frac{n \log(\pi)}{2} + \frac{|\Sigma|}{2} + \frac{1}{2} \text{tr}(\Sigma^{-1} \Sigma) \\ &= \frac{n \log(\pi)}{2} + \frac{|\Sigma|}{2} + \frac{n}{2} \end{aligned}$$

when  $n$  is the dimension of the multivariate gaussian in question.

- c) Adding a constant mean does not affect the differential entropy of a random variable. Therefore,  $H(\mathcal{N}(\mu, \Sigma)) = \frac{n \log(\pi)}{2} + \frac{|\Sigma|}{2} + \frac{n}{2}$ .

## 2 Problem 2

VAE implementation by me (KK Thuwajit / Kuroma). Code resources: my old github repository (<https://github.com/konkuad/GANime>) from several years ago.

```
[1]: #Imports
import numpy as np
import matplotlib.pyplot as plt
import torch
from torch import nn
import random
from tqdm import tqdm
import torchvision
import torchvision.datasets as datasets

#for consistency, all seeds are set to 69420
seed = 69420
random.seed(seed)
np.random.seed(seed)
torch.manual_seed(seed)
torch.cuda.manual_seed(seed)
```

### 2.1 2.1) VAE Encoder

This model attempts to transform real world data to  $\mathcal{N}(0, I)$ . It has three components: the convolutional encoder that maps images to a 512-dimensional vector, a fully connected layer for the predicted means, and another one for covariance. We use the reparameterization trick (i.e. <https://www.baeldung.com/cs/vae-reparameterization>) to randomly sample from latent space by shifting/scaling a randomly sampled standard normal vector by our computed mean and variances. The model keeps track of its KL-divergence to be backpropagated at training time.

```
[2]: class Encoder(nn.Module):
    def __init__(self, latent_size, img_channel):
        super(Encoder, self).__init__()

        self.encoder = nn.Sequential(
            nn.Conv2d(img_channel, 32, 4, 2, 1, bias=False),
            nn.BatchNorm2d(32),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Conv2d(32, 64, 4, 2, 1, bias=False),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(0.2, inplace=True),

            nn.Conv2d(64, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2, inplace=True),
```

```

        nn.Conv2d(128, 256, 4, 2, 1, bias=False),
        nn.BatchNorm2d(256),
        nn.LeakyReLU(0.2, inplace=True),
    )
    self.mu_fc = nn.Conv2d(256, latent_size, 2, 1, 0, bias=False)
    self.sigma_fc = nn.Conv2d(256, latent_size, 2, 1, 0, bias=False)
    self.kl = 0

    def forward(self, x, normal_generator):
        encoded = self.encoder(x)
        mu = self.mu_fc(encoded).mean([-2, -1]) #mean
        sigma = torch.exp(self.sigma_fc(encoded).mean([-2, -1])) #covariance
        z = mu + sigma * normal_generator.sample(mu.shape) #randomly sample from
        ↪ latent
        self.kl = (sigma**2 + mu**2 - torch.log(sigma) - 1/2).sum() # kl loss
        ↪ term
        return z

```

## 2.2 VAE Decoder

This convolutional neural network model maps latent vectors and reconstruct images from it. This decoder is objectively trained using the L2 reconstruction loss (though L1 would work under a different maximum likelihood assumption, as discussed in class).

```

[3]: class Decoder(nn.Module):

    def __init__(self, latent_size, img_channel):

        super(Decoder, self).__init__()

        self.conv_transpose_block_1 = nn.Sequential(
            nn.ConvTranspose2d(latent_size, 256, 4, 1, 0, bias=False),
            nn.BatchNorm2d(256),
            nn.LeakyReLU(0.2, inplace=True))

        self.conv_transpose_block_2 = nn.Sequential(
            nn.ConvTranspose2d(256, 128, 4, 2, 1, bias=False),
            nn.BatchNorm2d(128),
            nn.LeakyReLU(0.2, inplace=True))

        self.conv_transpose_block_3 = nn.Sequential(
            nn.ConvTranspose2d(128, 64, 4, 2, 1, bias=False),
            nn.BatchNorm2d(64),
            nn.LeakyReLU(0.2, inplace=True))

        self.conv_transpose_block_4 = nn.Sequential(
            nn.ConvTranspose2d(64, 32, 4, 2, 1, bias=False),

```

```

        nn.BatchNorm2d(32),
        nn.LeakyReLU(0.2, inplace=True))

    self.conv_transpose_block_5 = nn.Sequential(
        nn.ConvTranspose2d(32, img_channel, 1, 1, 0, bias=False),
        nn.Sigmoid())

    def forward(self, x):
        x = self.conv_transpose_block_1(x)
        x = self.conv_transpose_block_2(x)
        x = self.conv_transpose_block_3(x)
        x = self.conv_transpose_block_4(x)
        x = self.conv_transpose_block_5(x)
        return x

```

## 2.3 2.3) Dataloader

We load the MNIST dataset using torchvision's built-in loader. We normalize each pixel to (0,1) and resize to 32x32 for shape convenience.

```

[4]: transforms = torchvision.transforms.Compose([
    torchvision.transforms.Resize((32,32)),
    torchvision.transforms.ToTensor()
])

mnist_trainset = datasets.MNIST(root='./data', train=True, download=True,
    ↪transform=transforms)
data_loader = torch.utils.data.
    ↪DataLoader(mnist_trainset, batch_size=128, shuffle=True, num_workers=1)

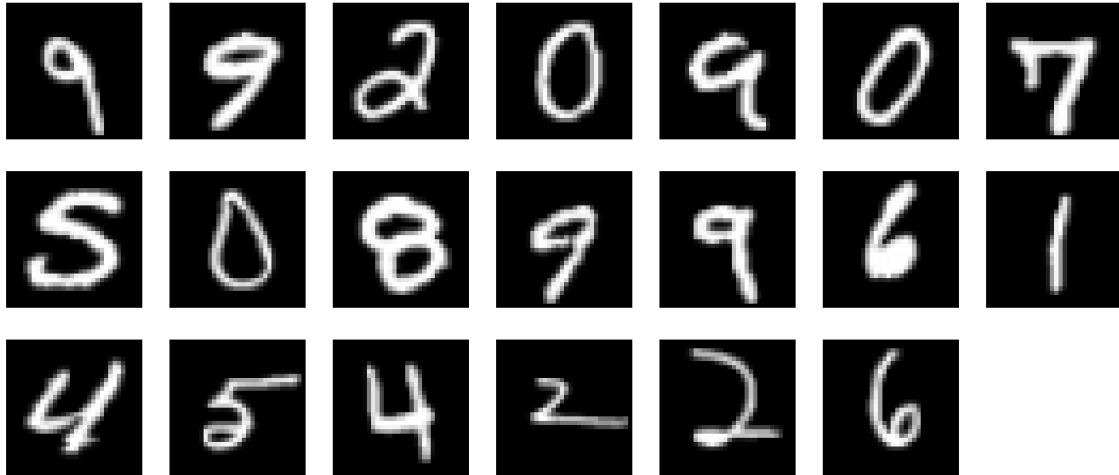
plt.gray()

def plotter(images):
    f = -(-len(images)//3)
    fig, ax = plt.subplots(3, f, figsize=(5*f, 15))
    for i in range(len(images)):
        ax[i%3, i//3].imshow(images[i, 0])
    for aa in ax:
        for aaa in aa:
            aaa.axis("off")
    plt.show()

for a, b in data_loader:
    plotter(a[0:20])
    break

```

<Figure size 640x480 with 0 Axes>



## 2.4 2.4) Training

We train the VAE with two losses: KL divergence and reconstruction (L2)

```
[5]: latent_size = 128
img_ch = 1
num_epochs = 20

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
normal_generator = torch.distributions.Normal(torch.tensor(0).float().
    ↪to(device), torch.tensor(1).float().to(device))

enc = Encoder(latent_size, img_ch).to(device)
dec = Decoder(latent_size, img_ch).to(device)

optimizer = torch.optim.Adam(list(enc.parameters())+list(dec.parameters()),
    ↪lr=0.0002, betas=(0.5, 0.999))
```

```
[6]: visualize_noise = torch.randn(16, latent_size, 1, 1).float().to(device)

for epoch in range(num_epochs):
    pbar = tqdm(enumerate(data_loader))
    enc.train()
    dec.train()
    count = 0
    kl_sum = 0
    l2_sum = 0
    for i, (data, _) in pbar:
        optimizer.zero_grad()
        x = data.to(device)
        b = x.shape[0]
```

```

z = enc(x, normal_generator)
kl_loss = enc.kl
reconstructed = dec(z.reshape(b,latent_size,1,1))
l2_loss = ((reconstructed - x)**2).sum()
vae_loss = kl_loss + l2_loss
vae_loss.backward()
optimizer.step()

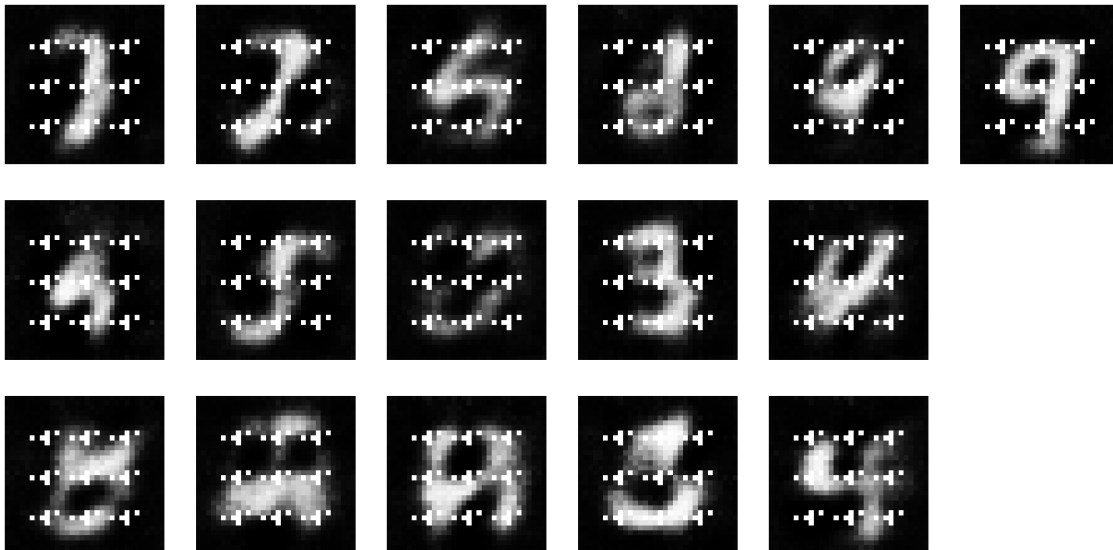
count += b
kl_sum += kl_loss.item()
l2_sum += l2_loss.item()

kl_loss_show = '{:.4f}'.format(kl_sum/count)
l2_loss_show = '{:.4f}'.format(l2_sum/count)
pbar.set_description(f'Iteration {i+1}/{len(data_loader)}\t[Epoch_
↪{epoch+1}/{num_epochs}]\tLosses:\tKL = {kl_loss_show} \tL2 = {l2_loss_show}')

# visualize every 4 epochs
if epoch%4==0:
    with torch.no_grad():
        dec.eval()
        pred = dec(visualize_noise.to(device).float())
        plotter(pred.cpu())

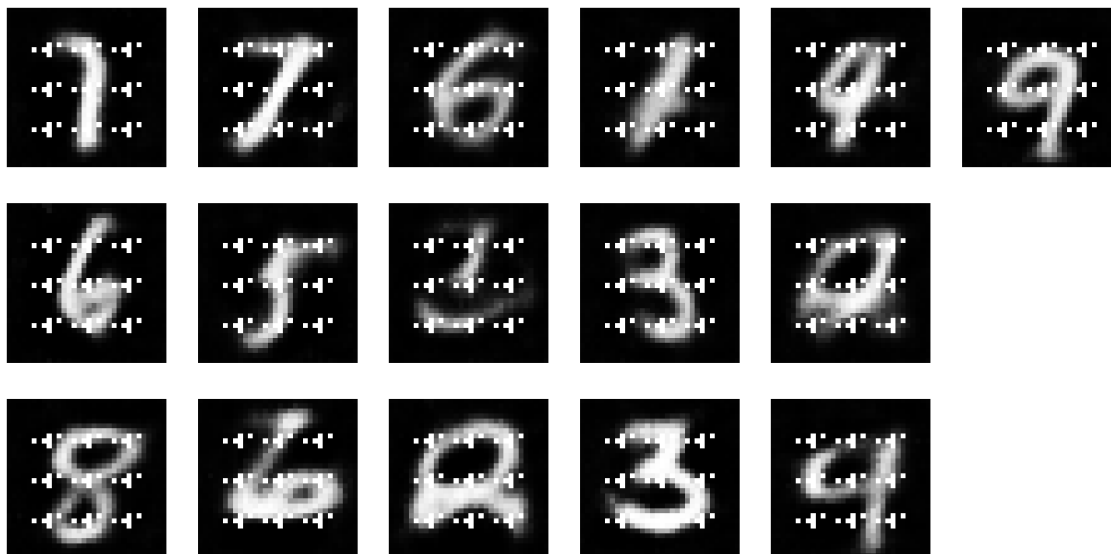
```

Iteration 469/469 [Epoch 1/20] Losses: KL = 54.6165 L2 = 74.8870: :  
469it [00:16, 28.60it/s]

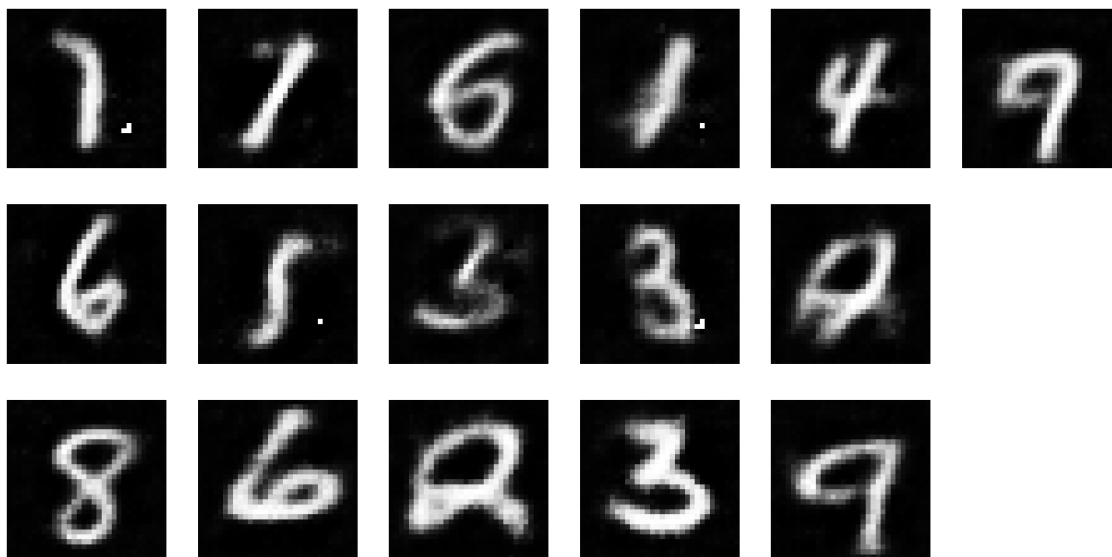


Iteration 469/469 [Epoch 2/20] Losses: KL = 55.1004 L2 = 59.6094: :  
469it [00:06, 69.10it/s]

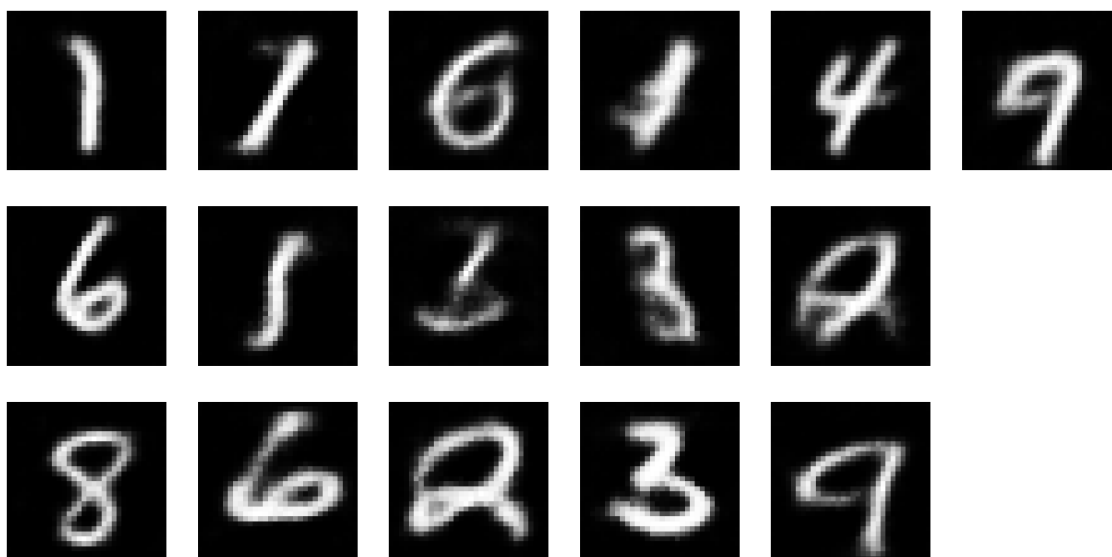
Iteration 469/469 [Epoch 3/20] Losses: KL = 55.4111 L2 = 57.1922: :  
 469it [00:07, 64.88it/s]  
 Iteration 469/469 [Epoch 4/20] Losses: KL = 55.5908 L2 = 55.7527: :  
 469it [00:07, 65.09it/s]  
 Iteration 469/469 [Epoch 5/20] Losses: KL = 55.7676 L2 = 54.7117: :  
 469it [00:07, 63.80it/s]



Iteration 469/469 [Epoch 6/20] Losses: KL = 55.8834 L2 = 53.9814: :  
 469it [00:07, 64.22it/s]  
 Iteration 469/469 [Epoch 7/20] Losses: KL = 55.9864 L2 = 53.4166: :  
 469it [00:07, 61.57it/s]  
 Iteration 469/469 [Epoch 8/20] Losses: KL = 56.0832 L2 = 51.3725: :  
 469it [00:06, 68.87it/s]  
 Iteration 469/469 [Epoch 9/20] Losses: KL = 57.0602 L2 = 19.0327: :  
 469it [00:06, 74.65it/s]



Iteration 469/469 [Epoch 10/20] Losses: KL = 56.9860 L2 = 16.4705: :  
 469it [00:06, 68.95it/s]  
 Iteration 469/469 [Epoch 11/20] Losses: KL = 57.0425 L2 = 15.7920: :  
 469it [00:07, 66.48it/s]  
 Iteration 469/469 [Epoch 12/20] Losses: KL = 57.0736 L2 = 15.4482: :  
 469it [00:06, 68.90it/s]  
 Iteration 469/469 [Epoch 13/20] Losses: KL = 57.1086 L2 = 15.2128: :  
 469it [00:07, 65.75it/s]



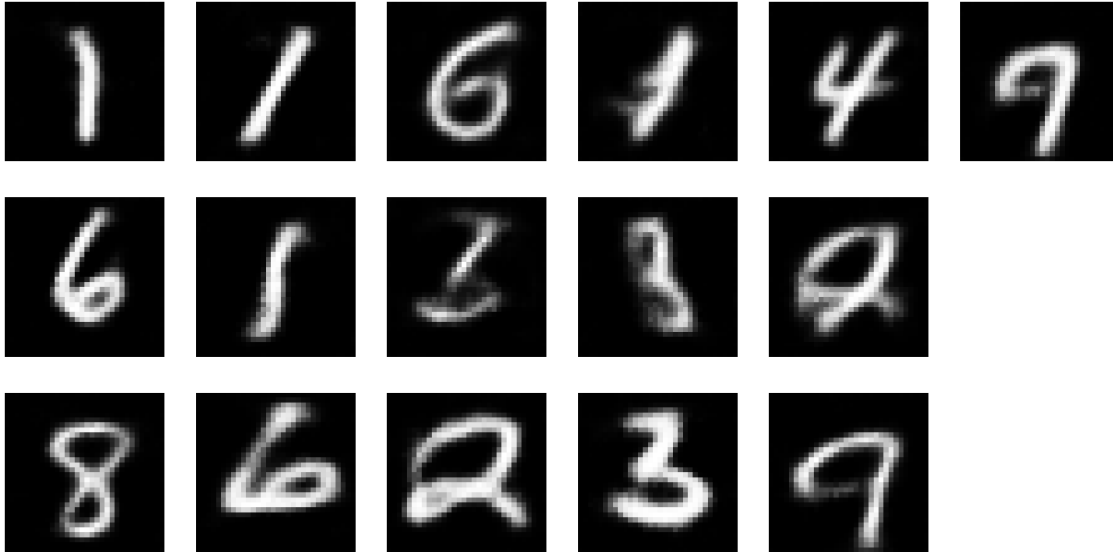
Iteration 469/469 [Epoch 14/20] Losses: KL = 57.1252 L2 = 14.9743: :



```

469it [00:06, 67.81it/s]
Iteration 469/469      [Epoch 15/20]   Losses: KL = 57.1699    L2 = 14.7871: :
469it [00:07, 63.17it/s]
Iteration 469/469      [Epoch 16/20]   Losses: KL = 57.1752    L2 = 14.6092: :
469it [00:07, 65.41it/s]
Iteration 469/469      [Epoch 17/20]   Losses: KL = 57.1972    L2 = 14.4964: :
469it [00:07, 66.84it/s]

```



```

Iteration 469/469      [Epoch 18/20]   Losses: KL = 57.2064    L2 = 14.3362: :
469it [00:07, 64.35it/s]
Iteration 469/469      [Epoch 19/20]   Losses: KL = 57.2227    L2 = 14.1998: :
469it [00:07, 66.23it/s]
Iteration 469/469      [Epoch 20/20]   Losses: KL = 57.2319    L2 = 14.0880: :
469it [00:07, 65.52it/s]

```

## 2.5 2.5) Random generation

We randomly generate images using our finished VAE. Observation: VAE's use of L2 to approximate its variational lower bound causes its generations to possess soft, undefined edges (unlike GANs).

```

[7]: with torch.no_grad():
    dec.eval()
    pred = dec(visualize_noise.to(device).float())
    plotter(pred.cpu())

```

